
MARKOV MODELLING

A PREPRINT

Kshitij Bichave*

March 3, 2021

Keywords Reinforcement Learning · Markov-chain · Markov-Reward · Markov Process · Return · Value Function · Return · Discount Factor · Markov Decision Process · MDP Control

Markov Process/Chain:

Sequence of random states with Markov property (that is, $p(s_{t+1}|s_t, a_t) = p(s_{t+1}|h_t, a_t)$)

Definition of Markov Process:

- S is a set of states, such that $s \in S$
- P is a dynamics model which specifies $P(s_{t+1} = s' | s_t = s)$

Discount Factor:

Used to avoid infinite return so that expected sum of return is bounded. If $\gamma = 0$, agent only cares about immediate reward. If $\gamma = 1$, future reward is as beneficial as immediate reward. γ can be set to 1 if the horizon is finite

Markov Reward Process (MRP):

Markov reward process is a Markov chain with addition/inclusion of reward term. Because MRP includes rewards and rewards can be weighted for how much we care about the future rewards, γ is included in MRP.

Definition of Markov Process:

- S is a set of states, such that $s \in S$
- P is a dynamics model which specifies $P(s_{t+1} = s' | s_t = s)$
- R is a reward function $R(s_t = s) = \mathbb{E}[r_t | s_t = s]$
- Discount factor $\gamma \in [0, 1]$

Return & Value Function:

Horizon: Number of time steps in each episode, can be infinite or finite.

Return G_t (for MRP) is a discounted sum of rewards from time step t to horizon. $G_t = r_t + \gamma r_{t+1} + \gamma^2 r_{t+2} + \dots$

State Value Function $V(s)$ (for MRP) is the expected return from starting state s . $V(s) = \mathbb{E}[G_t | s_t = s] = \mathbb{E}[r_t + \gamma r_{t+1} + \gamma^2 r_{t+2} + \dots | s_t = s]$. The State value function can be deterministic or stochastic. If deterministic, the expected return is always the same, because there is only a single next state the agent can go to.

Computing Value for Markov Reward Process:

*GitHub: github.com/kbichave | Email: k.bichave@gmail.com | LinkedIn: [linkedin.com/in/kshitijbichave](https://www.linkedin.com/in/kshitijbichave)

Using simulation:

- Take starting state distribution and generate large number of episodes
- Average returns
- Concentration inequalities bound how quickly average concentrates to expected value

Advantage: Does not require assumption of Markov

Given Markov property:

$$V(s) = R(s) + \gamma \sum_{s' \in S} P(s'|s) V(s').$$

Here, $R(s)$ is the immediate reward and the second term in the equation is the sum of future rewards discounted by discount factor γ .

Using Dynamic Programming (iterative algorithm):

- Initialize $V_0(s) = 0$ for all s
- for all $k = 1$ until convergence (condition for convergence could be $|V_k - V_{k-1}| < \epsilon$, where ϵ is some allowed error term
 - For all $s \in S$, $V_k(s) = R(s) + \gamma \sum_{s' \in S} P(s'|s) V_{k-1}(s')$

Markov Decision Process (MDP):

Markov Decision Processes are Markov Reward Process with addition/inclusion of actions

Definition of Markov Decision Process:

- S is a set of states, such that $s \in S$
- A is a set of actions, such that $a \in A$
- P is a dynamics model for each action which specifies $P(s_{t+1} = s' | s_t = s, a_t = a)$
- R is a reward function $R(s_t = s, a_t = a) = \mathbb{E}[r_t | s_t = s, a_t = a]$
- Discount factor $\gamma \in [0, 1]$

MDP is defined as a tuple: (S, A, P, R, γ)

MDP Policies:

Policies specifies which action to take in each state. The policies can be deterministic or stochastic. That is, does the same action when in that state, or has a distribution over actions and thus may land in different states.

Policy: $\pi(a|s) = P(a_t = a | s_t = s)$

MDP + Policy:

Markov Reward Process, MRP $(S, R^\pi, P^\pi, \gamma)$

$$R^\pi(s) = \sum_{a \in A} \pi(a|s) R(s, a)$$

$$P^\pi(s'|s) = \sum_{a \in A} \pi(a|s) P(s'|s, a)$$

This implies that we have fixed policies, simulation, analytic solution or dynamic solution could be used to compute value of the policy.

Dynamic programming now becomes:

- Initialize $V_0(s) = 0$ for all s
- for all $k = 1$ until convergence (condition for convergence could be $|V_k - V_{k-1}| < \epsilon$, where ϵ is some allowed error term
 - For all $s \in S$, $V_k^\pi(s) = r(s, \pi(a|s)) + \gamma \sum_{s' \in S} P(s'|s, \pi(a|s)) V_{k-1}^\pi(s')$

This is also called Bellman backup for a particular policy.

MDP Control:

MDP control now talks about not just evaluating policies but, learning, that is computing the optimal policy, $\pi^*(a|s) = \operatorname{argmax}_a V^\pi(s)$.

There exists a unique optimal value function and a unique optimal policy for a MDP.

In an infinite horizon problem optimal policy is deterministic and stationary (does not depend on time step).

Given 7 discrete states and 2 actions, the number of deterministic policies for this example are 2^7 .

Optimal policy is not always unique. But value function is unique.

MDP Policy Iteration:

In policy iteration, you keep track of what the optimal policy might be, we evaluate its value and then improve it.

- set $i = 0$
- Initialize π_0 randomly for all states s
- While $i == 0$ or $\|\pi_i - \pi_{i-1}\| > 0$ (L1-norm, measures if the policy changed for any state):
 - $V^{\pi_i} \leftarrow$ MDP V function policy evaluation of π_i
 - $\pi_{i+1} \leftarrow$ Policy improvement
 - $i = i + 1$

In order to define how, policy could improve, define state-action value $Q^\pi(s, a) = R(s, a) + \gamma \sum_{s' \in S} P(s'|s, a) V^\pi(s')$. That is take action a , then follow the policy π .

Policy Improvement:

After we have a policy from policy evaluation, and thus we have value of it. That is, we know V^{π_i} . We also know what the dynamics and reward model is. So now,

Compute the state-action value of policy π_i

For $s \in S$ and $a \in A$:

$$Q^{\pi_i}(s, a) = R(s, a) + \gamma \sum_{s' \in S} P(s'|s, a) V^{\pi_i}(s')$$

Compute new policy π_{i+1} for all $s \in S$. That is choose a policy such that it maximizes Q^{π_i} .

$$\pi_{i+1}(s) = \operatorname{argmax}_a Q^{\pi_i}(s, a) \quad \forall s \in S$$

In theory, $\operatorname{argmax}_a Q^{\pi_i}(s, a) \geq Q^{\pi_i}(s, \pi_i(s))$. Because, $a = \pi_i(s)$ and thus will be same or better according to argmax . This is guaranteed because of monotonic improvement of policy. That is, something is monotonic is new is equal or better than older. That is $V^{\pi_{i+1}} \geq V^{\pi_i}$ with strict inequality if π_i is sub-optimal, where π_{i+1} is the new policy we get from improvement on π_i . So, if $\pi_{i+1} = \pi_i$ for some i , then the policy can never get better beyond π_i .

Bellman Equation and Bellman Backup Operator

$$V^\pi(s) = R(s) + \gamma \sum_{s' \in S} P(s'|s) V^\pi(s') \leftarrow \text{Bellman equation}$$

Can be thought of Bellman backup operator. That is when applied to value function, return a new value function and improves the values if possible. BV yields a value function over all states s .

$$BV(s) = \max_a R(s, a) + \gamma \sum_{s' \in S} p(s'|s, a) V(s')$$

Value Iteration:

Idea is that maintain optimal value of starting in a state s if have a finite number of steps k left in episode. Iterate to consider longer and longer episodes. Value iteration assumes you have optimal value and policy only if you are going to act for k time steps (finite horizon).

- set $k = 1$
- Initialize $V_0(s)$ randomly for all states s
- Loop until convergence:
 - For each state s ,

$$V_{k+1}(s) = \max_a R(s, a) + \gamma \sum_{s' \in S} p(s'|s, a) V_k(s').$$

That is Bellman backup on value function, $V_{k+1} = BV_k$.

$$- \pi_{i+1}(s) \operatorname{argmax}_a R(s, a) + \gamma \sum_{s' \in S} P(s'|s, a) V_k(s')$$

Bellman operator can be initialized by fixing the policy B^π , instead of taking \max_a . So to do policy evaluation, repeatedly apply operator until V stops changing.

Value iteration will converge if discount factor $\gamma < 1$ or ends up in terminal state with probability 1. Therefore, it can be said that Bellman backup is contraction if $\gamma < 1$. If applied to two different value functions, distance between value functions shrinks after applying Bellman equation to each.